

# CSE 421/521

# Introduction to Operating Systems

Farshad Ghanei

## Lecture - 05

## Project-1 Discussion

\* Slides adopted from Prof Kosar and Dantu at UB, "Operating System Concepts" book and supplementary material by A. Silberschatz, P.B. Galvin, and G. Gagne. Wiley Publishers

 University at Buffalo  
Department of Computer Science  
and Engineering  
School of Engineering and Applied Sciences

*The materials provided by the instructor in this course are for the use of the students enrolled in the course only.  
Copyrighted course materials may not be further disseminated without instructor permission.*



# Summary

- Threads
  - Concurrent programming
  - Why threads?
  - Threads vs Processes
  - Threading examples
  - Thread pools
  - Threading implementation & multithreading models
  - Threading issues
    - Semantics of `fork()` and `exec()`
    - Thread cancellation
    - Signal handling



# Today

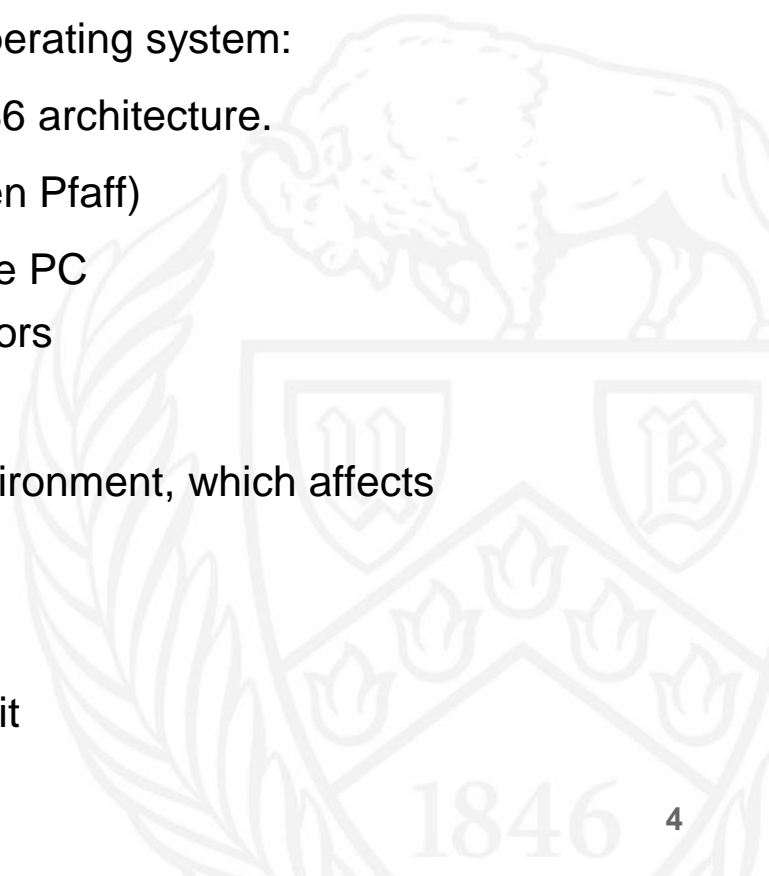
- Pintos projects:
  - Project-1: Threads
    - Step 1: Preparation
    - Step 2: Setting Up Pintos
    - Step 3: Design Document
    - Step 4: Implementation
    - Step 5: Testing
  - Project-2: User Programs
  - Project-3: Virtual Memory
  - Project-4: File Systems



# Pintos

Your programming assignments are based on Pintos operating system:

- A Simple operating system framework for the 80x86 architecture.
- Developed by Stanford University (Originally by Ben Pfaff)
- Could theoretically run on a regular IBM-compatible PC
- Practically, it runs using **Bochs** and **QEMU** simulators  
(OS in OS ?)
- Each and every student might have a different environment, which affects development, and running
- We provide a VirtualBox image with Ubuntu 16.04
  - Download VirtualBox and the image, and run it
- We'll run Bochs and QEMU within the VM image.  
(OS in OS in OS ?)

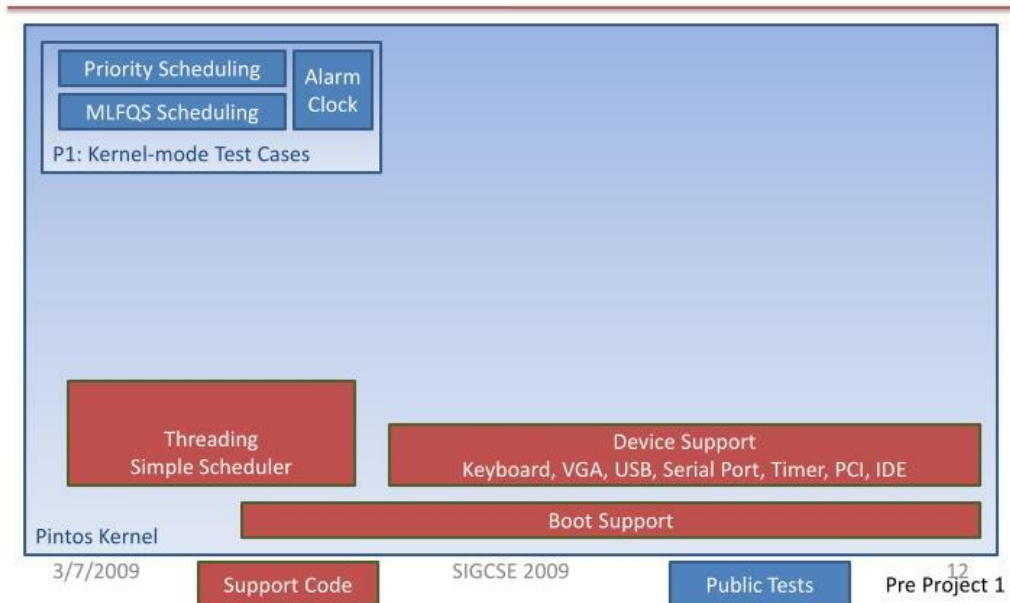


# Pintos Projects

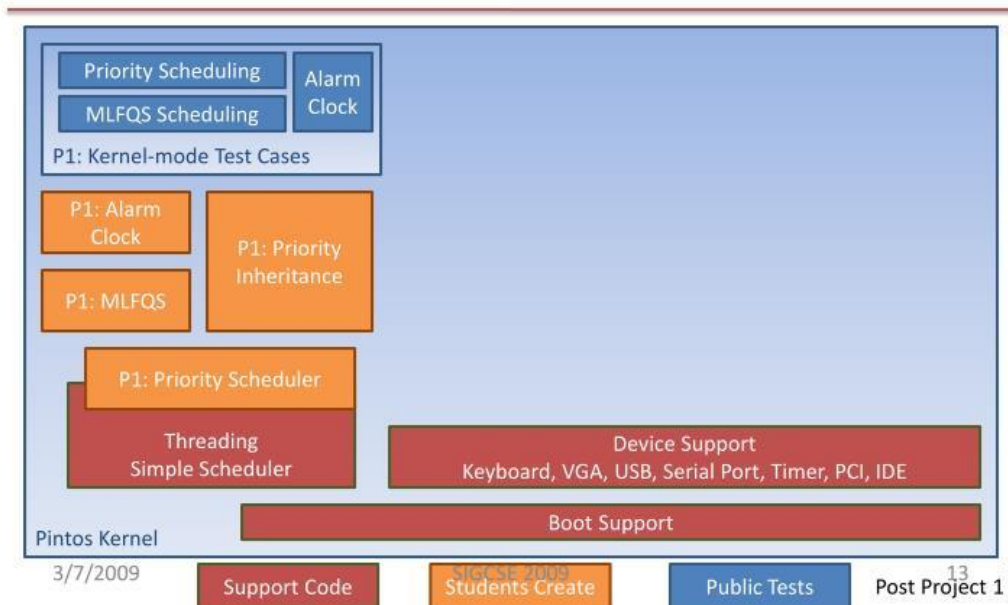
- Threads (CSE 421/521 Project-1)
- User Programs (CSE 421/521 Project-2)
- Virtual Memory
- File Systems



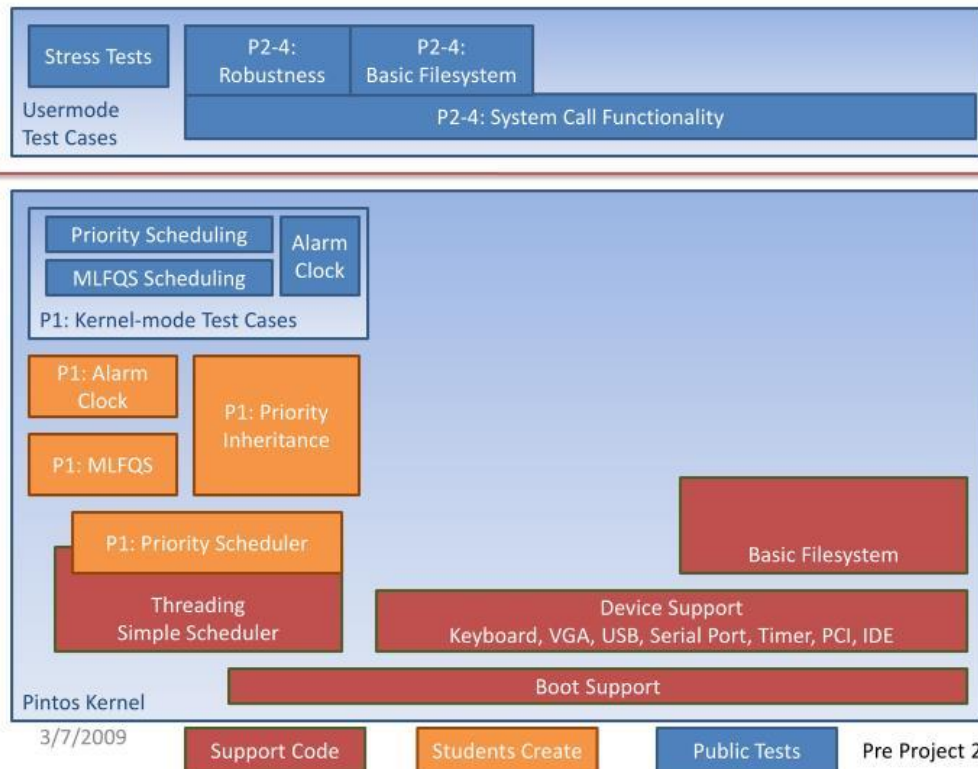
# Pintos (1/8) - Pre Project 1



# Pintos (2/8) - Post Project 1

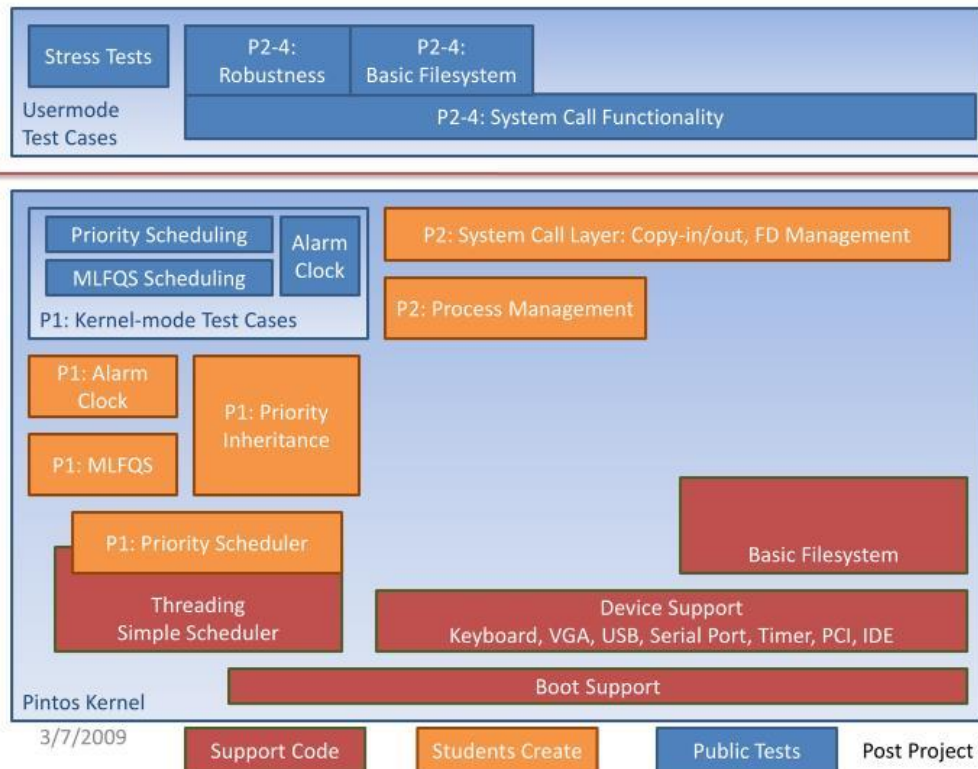


# Pintos (3/8) - Pre Project 2

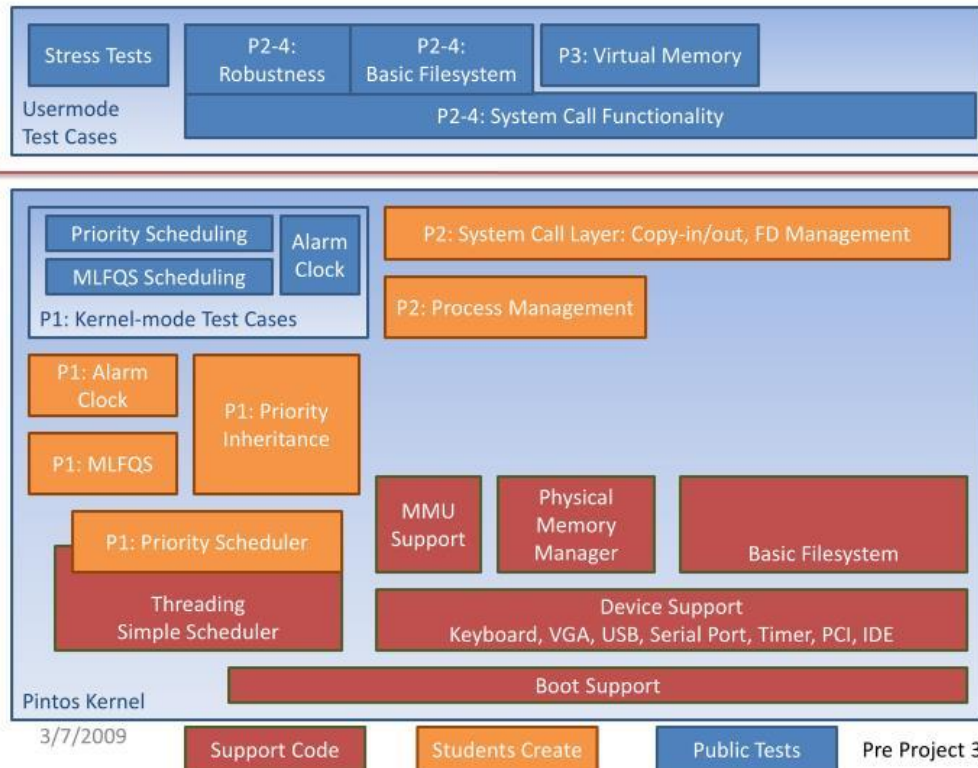




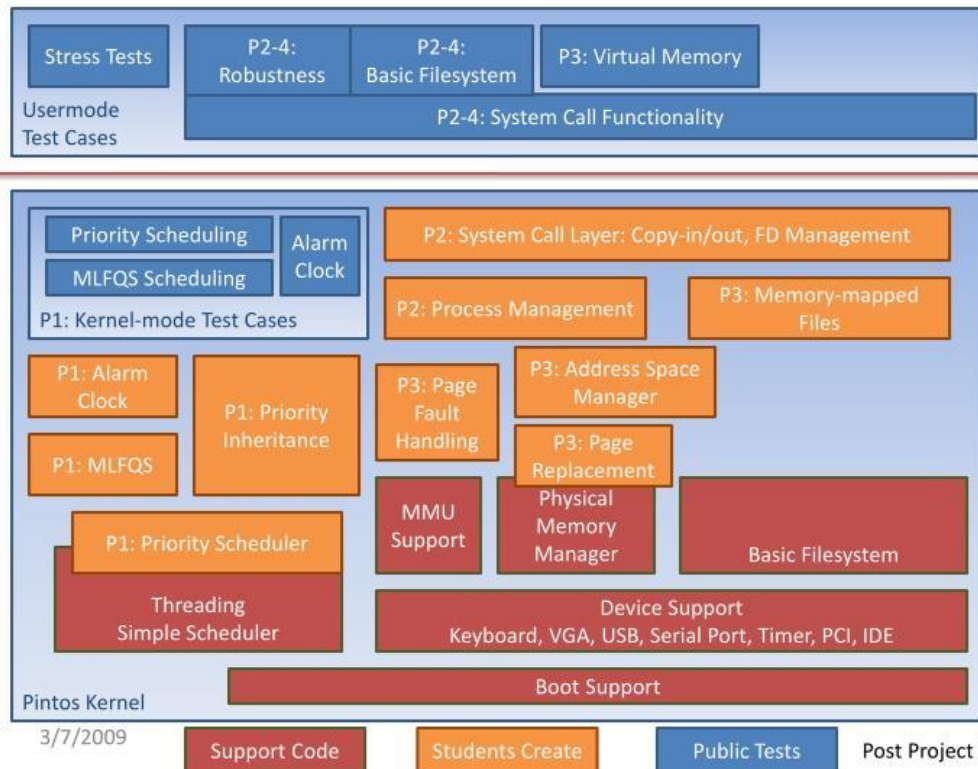
# Pintos (4/8) - Post Project 2



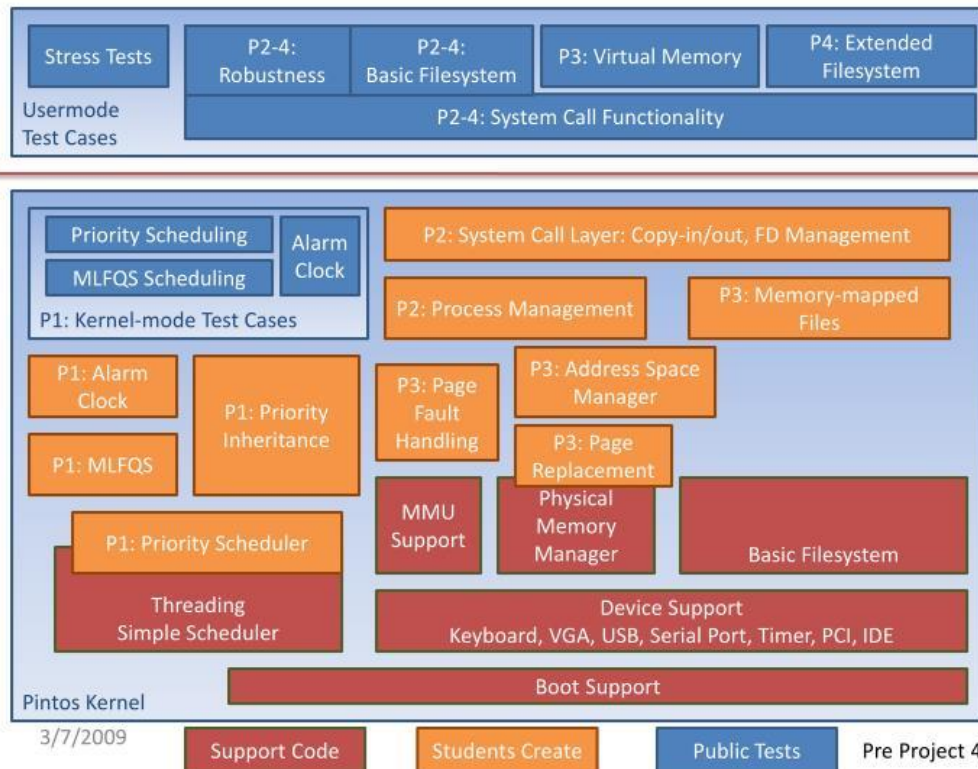
# Pintos (5/8) - Pre Project 3



# Pintos (6/8) - Post Project 3



# Pintos (7/8) - Pre Project 4



3/7/2009

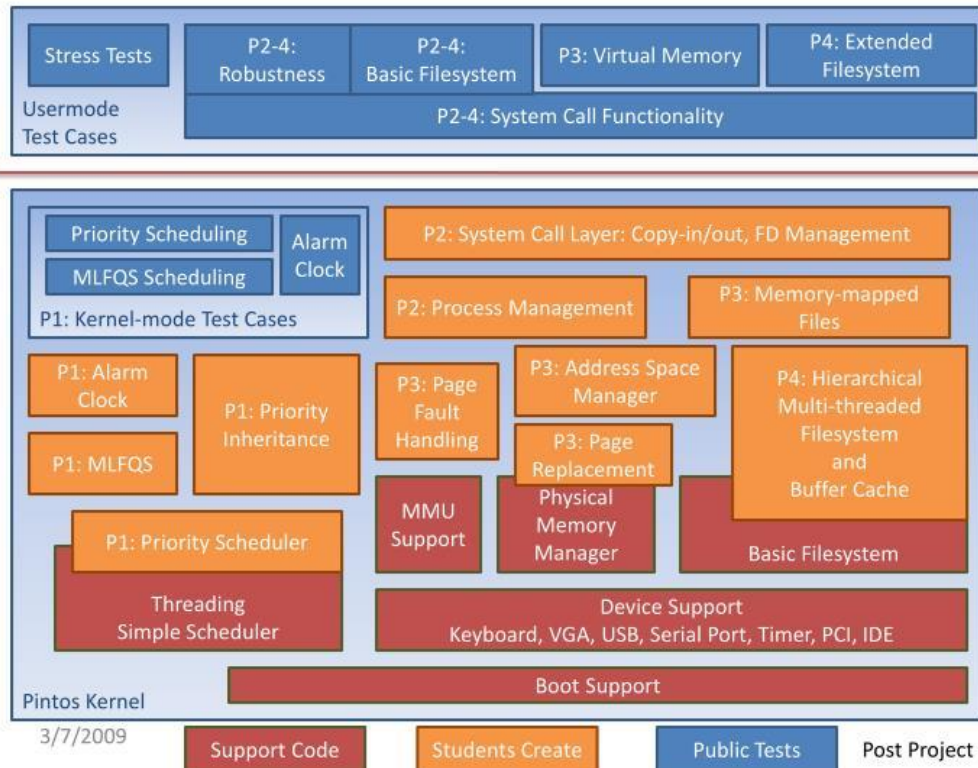
Support Code

Students Create

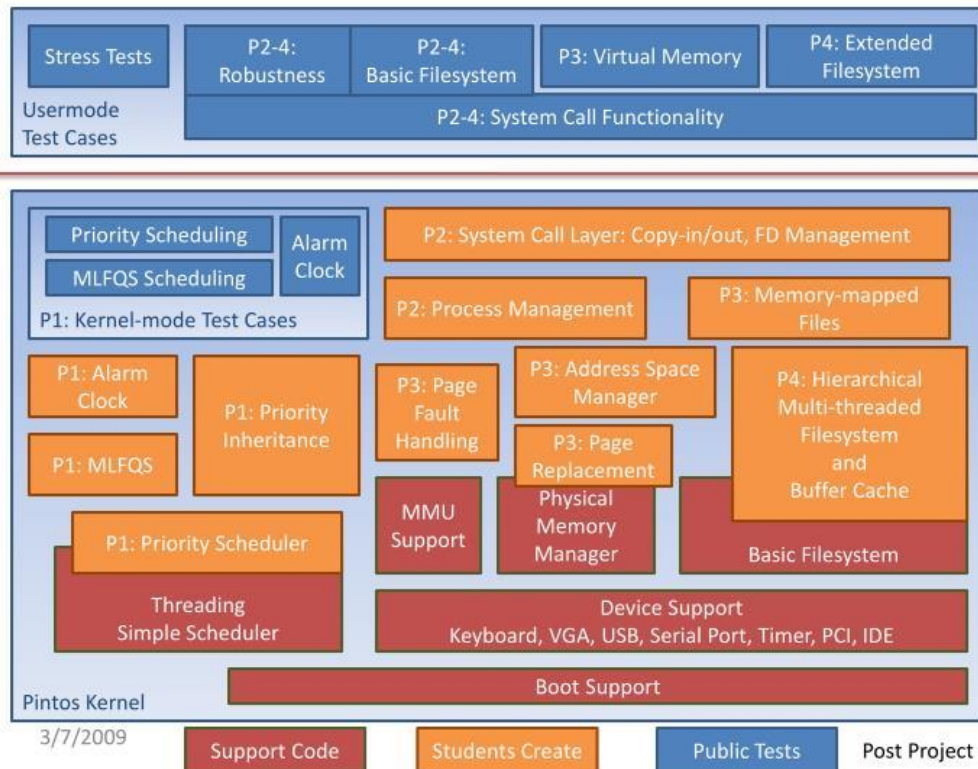
Public Tests

Pre Project 4

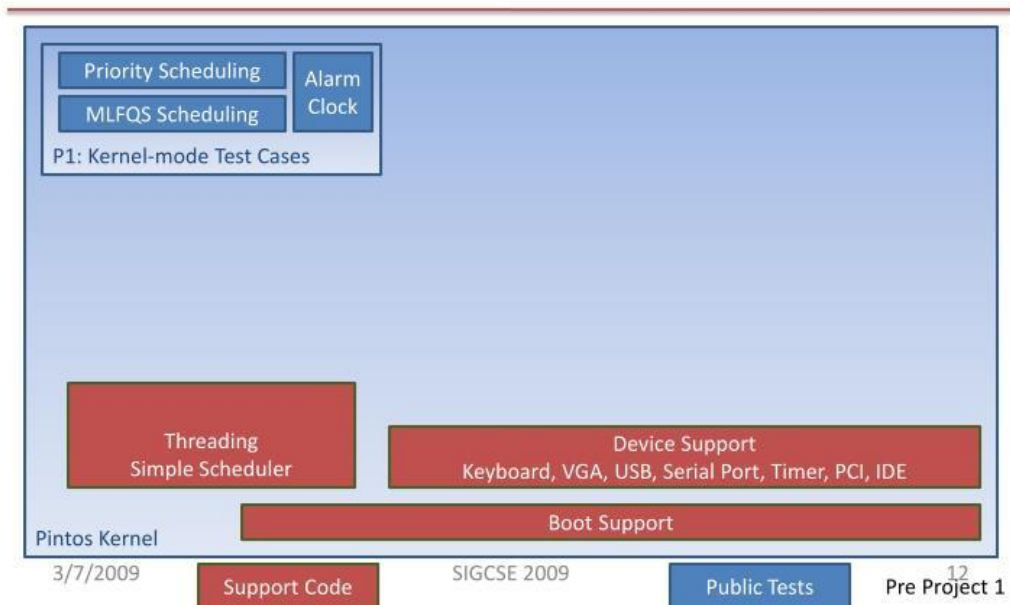
# Pintos (8/8) - Post Project 4



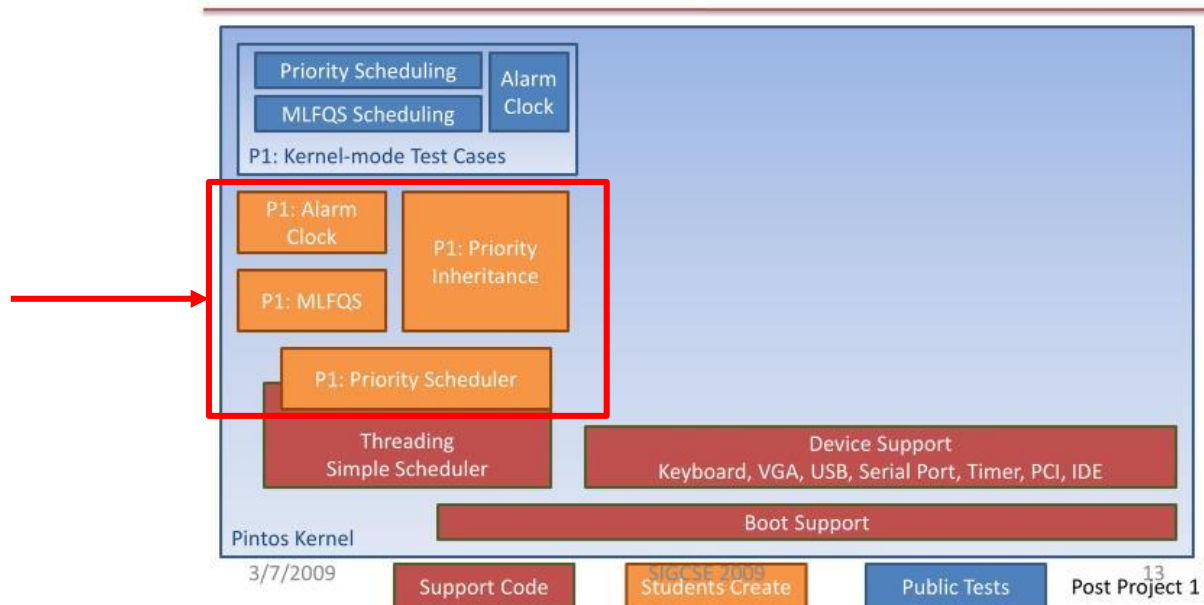
# Pintos - After full implementation (Post Project 4)



# Pintos - You Start from Here (Pre Project 1)

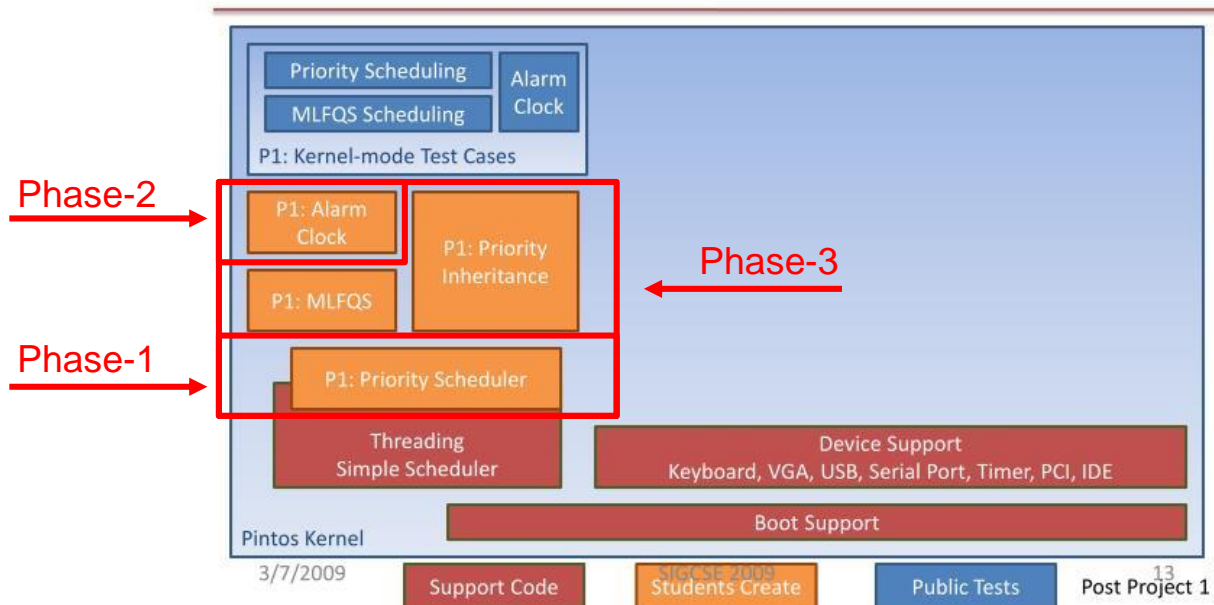


# Pintos - You Will Implement This (Post Project 1)





# Pintos - You Will Implement This (Post Project 1)



# Project-1: Threads

- Step 1: Preparation
- Step 2: Setting Up Pintos
- Step 3: Design Document
- Step 4: Implementation
- Step 5: Testing




# Step 1: Preparation

Readings from zyBooks:

- Chapters 2, 3, 4

Readings from Pintos documentation:

- Chapter 1 - Introduction
- Chapter 2 - Project 1: Threads
- Appendix A - Reference Guide [A.1, A.2, A.3, A.4]
- Appendix B - 4.4BSD Scheduler
- Appendix C - Coding Standards
- Appendix D - Project Documentation



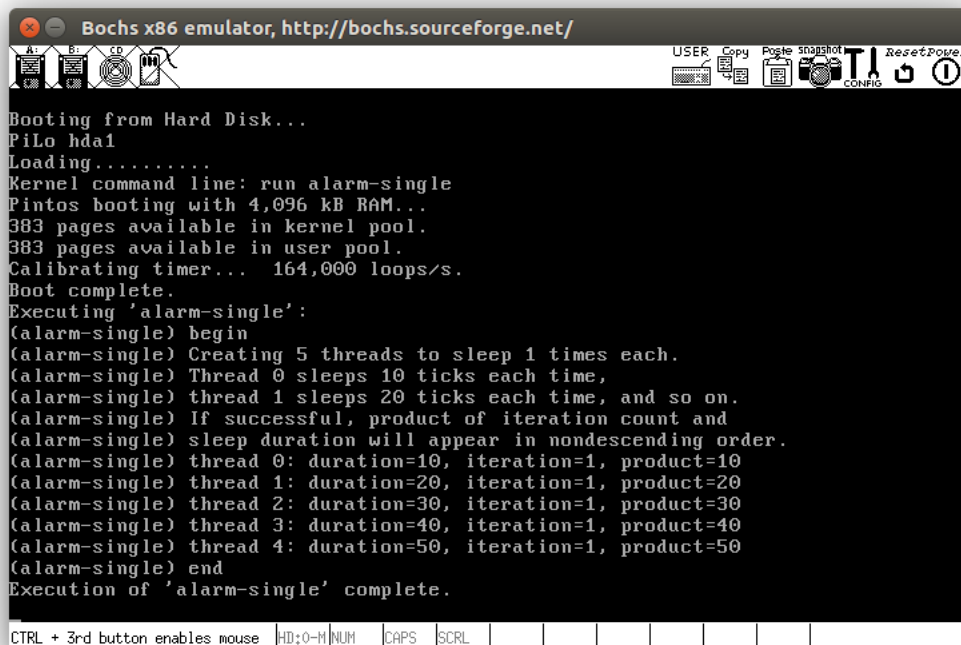
## Step 2: Setting Up Pintos

- Use the **VM** we have prepared for you
- Verify setup
  - Compile:
    - `cd $PINTOSDIR/src/threads`
    - `make`
  - Test:
    - `cd build`
    - `pintos run alarm-single`

Keep this environment variable valid



## Step 2: Setting Up Pintos



```

Bochs x86 emulator, http://bochs.sourceforge.net/
USER Copy Paste snapshot CONFIG Reset/Power

Booting from Hard Disk...
PiLo hda1
Loading.....
Kernel command line: run alarm-single
Pintos booting with 4,096 kB RAM...
383 pages available in kernel pool.
383 pages available in user pool.
Calibrating timer... 164,000 loops/s.
Boot complete.
Executing 'alarm-single':
(alarm-single) begin
(alarm-single) Creating 5 threads to sleep 1 times each.
(alarm-single) Thread 0 sleeps 10 ticks each time,
(alarm-single) thread 1 sleeps 20 ticks each time, and so on.
(alarm-single) If successful, product of iteration count and
(alarm-single) sleep duration will appear in nondescending order.
(alarm-single) thread 0: duration=10, iteration=1, product=10
(alarm-single) thread 1: duration=20, iteration=1, product=20
(alarm-single) thread 2: duration=30, iteration=1, product=30
(alarm-single) thread 3: duration=40, iteration=1, product=40
(alarm-single) thread 4: duration=50, iteration=1, product=50
(alarm-single) end
Execution of 'alarm-single' complete.

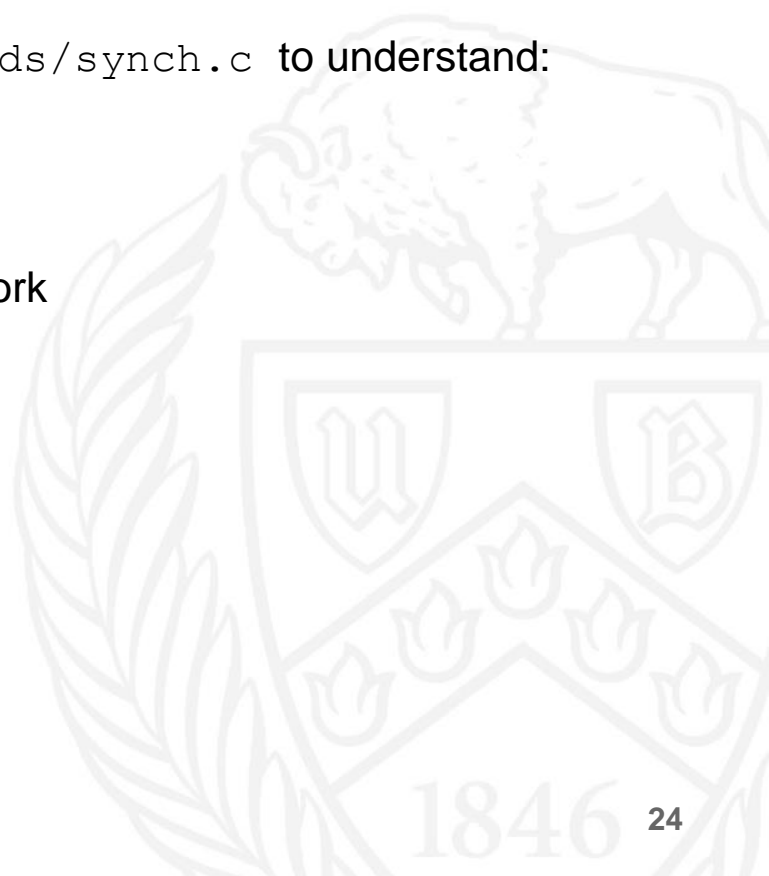
CTRL + 3rd button enables mouse HD:0-H NUM CAPS SCRL
    
```

# Get Familiar with the Code

- The first task is to read and understand the code for the initial thread system
  - Under the `src/threads/` directory
- Pintos already implements thread creation and thread completion, a simple scheduler to switch between threads, and synchronization primitives (semaphores, locks, condition variables, and optimization barriers).
- For a brief overview of the files in the `src/threads/` directory, please see **Section 2.1.2 Source Files** in the Pintos documentation.

# Pintos Thread System

- Read `src/threads/thread.c` and `src/threads/synch.c` to understand:
  - How the switching between threads occur
  - How the provided scheduler works
  - How the various synchronization primitives work
- These are **critical** steps.



# Important Directories

- `src/threads/`

Source code for the base kernel, which you will modify starting in project-1.

- `src/devices/`

Source code for I/O device interfacing: keyboard, timer, disk, etc. You will modify the timer implementation in project-1. Otherwise you should have no need to change this code.

- `src/lib/kernel/`

Parts of the C library that are included only in the Pintos kernel. Feel free to reuse this code.

- `src/tests/`

Tests for each project. You can read and modify this code to better understand your implementation. However, we will replace them with originals before we run tests.



# Files of Interest

- `thread.c` and `thread.h`

Basic thread support. Most of project-1 work.

- `synch.c` and `synch.h`

Synchronization primitives which you can use in all projects

- `devices/timer.c` and `devices/timer.h`

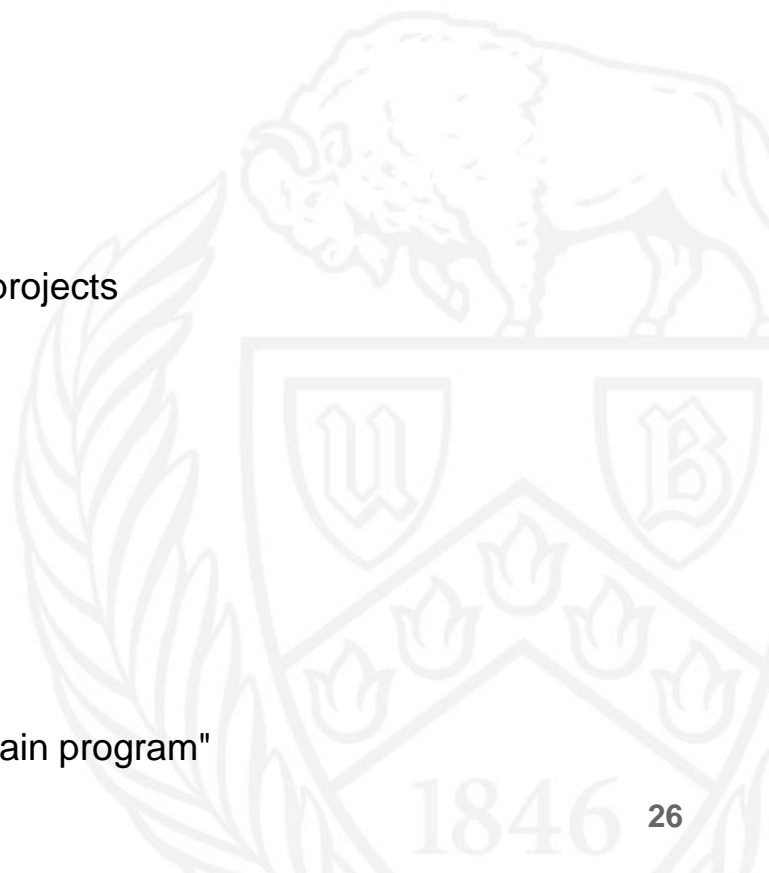
Timer ticks, has to be modified for project-1

- `lib/kernel/list.c`

Linked list implementation, feel free to reuse

- `init.c` and `init.h`

Kernel initialization, including `main()`, the kernel's "main program"



# Pintos Thread System

```

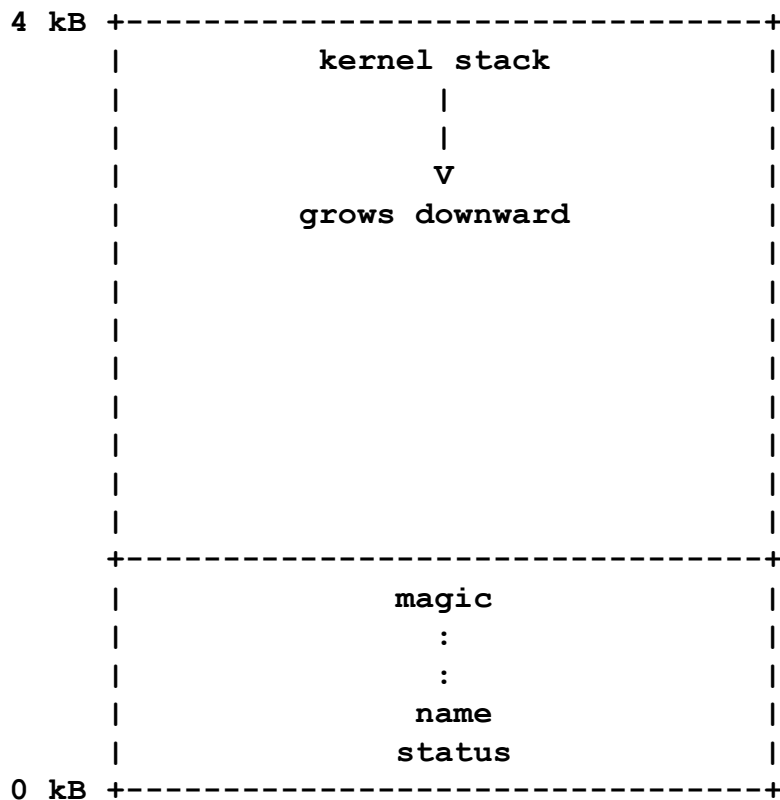
struct thread
{
    /* Owned by thread.c. */
    tid_t tid;                /* Thread identifier. */
    enum thread_status status; /* Thread state. */
    char name[16];            /* Name (for debugging purposes). */
    uint8_t *stack;           /* Saved stack pointer. */
    int priority;             /* Priority. */
    struct list_elem allelem; /* List element for all threads list. */

    /* Shared between thread.c and synch.c. */
    struct list_elem elem;    /* List element. */

#ifdef USERPROG
    /* Owned by userprog/process.c. */
    uint32_t *pagedir;        /* Page directory. */
#endif

    /* Owned by thread.c. */
    unsigned magic;           /* Detects stack overflow. */
};
    
```

# Pintos Thread System



## Step 3: Design Document

Use the template in `doc/` directory:

- `threads.tpl`
  - `userprog.tpl`
  - `vm.tpl`
  - `filesystem.tpl`
- 
- Copy the `threads.tpl` file for your design doc submission.

## Step 3: Design Document

```
+-----+
|           CS 140           |
| PROJECT 1: THREADS |
|   DESIGN DOCUMENT   |
+-----+
```

---- GROUP ----

>> Fill in the names and email addresses of your group members.

FirstName LastName <email@domain.example>

FirstName LastName <email@domain.example>

FirstName LastName <email@domain.example>

---- PRELIMINARIES ----

>> If you have any preliminary comments on your submission, notes for the

>> TAs, or extra credit, please give them here.

>> Please cite any offline or online sources you consulted while

>> preparing your submission, other than the Pintos documentation, course

>> text, lecture notes, and course staff.

# Step 3: Design Document

ALARM CLOCK

=====

---- DATA STRUCTURES ----

>> A1: Copy here the declaration of each new or changed `struct' or  
>> `struct' member, global or static variable, `typedef', or  
>> enumeration. Identify the purpose of each in 25 words or less.

---- ALGORITHMS ----

>> A2: Briefly describe what happens in a call to timer\_sleep(),  
>> including the effects of the timer interrupt handler.

>> A3: What steps are taken to minimize the amount of time spent in  
>> the timer interrupt handler?

---- SYNCHRONIZATION ----

>> A4: How are race conditions avoided when multiple threads call  
>> timer\_sleep() simultaneously?

>> A5: How are race conditions avoided when a timer interrupt occurs  
>> during a call to timer\_sleep()?

# Step 4: Implementation

## Phase-1

1. Priority Scheduling

## Phase-2

1. Alarm Clock

## Phase-3

1. Priority Scheduling (cont.)
2. Priority Donation
3. Multilevel Feedback Queue Scheduler (MLFQS)

Effort and time needed: Phase-3 > Phase-2 > Phase-1



# Task 1: Implement Priority Scheduler

- Ready thread with highest priority gets the processor
- When a thread is added to the ready list that has a higher priority than the currently running thread, immediately yield the processor to the new thread
- Implementation details:
  - Compare priority of the thread being added to the ready list with that of the running thread (preemptive)
  - Select next thread to run based on priorities

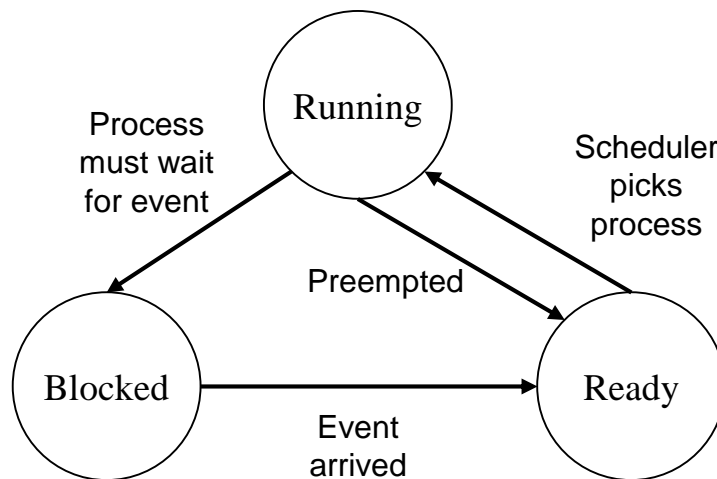


# Task 1: Implement Priority Scheduler

- Use `thread_yield()` to implement preemption.
  - Current thread (“Running”) is moved to READY state, added to READY list.
  - Then scheduler is invoked. Picks a new READY thread from READY list.
  - Case 1: There is only 1 READY thread. Thread is rescheduled right away.
  - Case 2: There are other READY threads.
    - 2.a) Another thread has higher priority - It is scheduled
    - 2.b) Another thread has same priority - It is scheduled provided the previously running thread was inserted in tail of ready list.
    - 2.c) Other threads have lower priority - Current thread gets rescheduled

# Task 1: Implement Priority Scheduler

- `thread_yield()` is a call you can use whenever you identify a need to preempt current thread.
- **Exception:** inside an interrupt handler, use `intr_yield_on_return()` instead.



Re-implement `next_thread_to_run()` for priority scheduling

# Task 1: Implement Priority Scheduler

- One or more threads are in `ready_list`, and getting scheduled by the scheduler.
- Where should we intervene (i.e. preempt current thread)?
  - Change to current threads. E.g. if priority of a thread changes
    - `void thread_set_priority (int new_priority)`
  - A new thread with higher priority. I.e. creation of a new thread.
    - `tid_t thread_create (const char *, int, thread_func *, void *)`

## Task 2: Implement Alarm Clock

- Reimplement `timer_sleep()` in `devices/timer.c` without busy-waiting

```
/* Sleeps for approximately TICKS timer ticks.  Interrupts must  
   be turned on. */
```

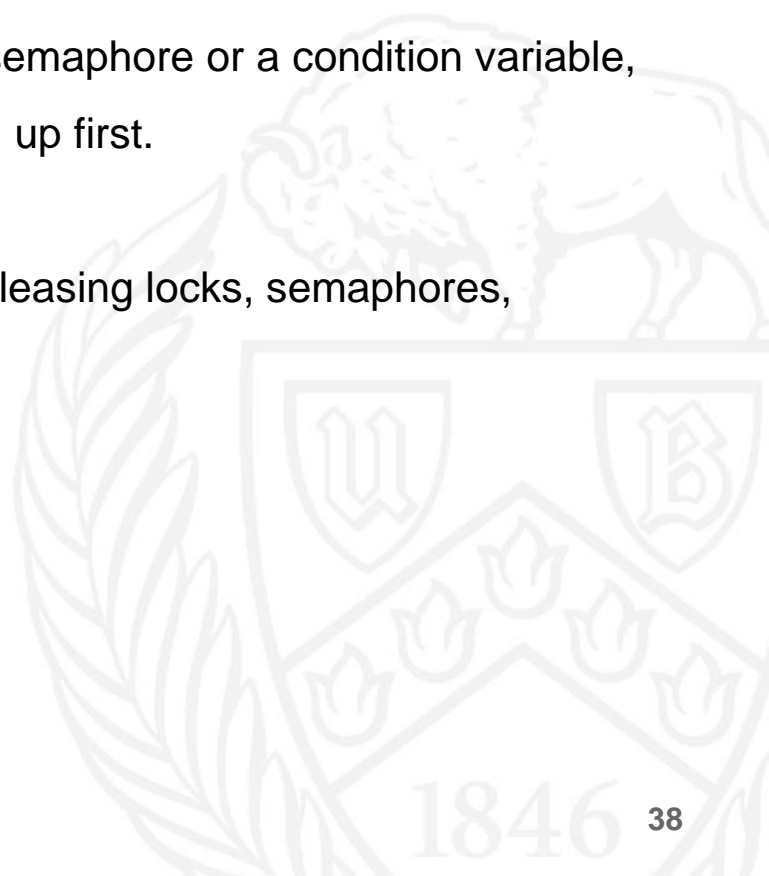
```
void  
timer_sleep (int64_t ticks)  
{  
    int64_t start = timer_ticks ();  
  
    ASSERT (intr_get_level () == INTR_ON);  
    while (timer_elapsed (start) < ticks)  
        thread_yield ();  
}
```

- Implementation details:
  - Remove thread from ready list and put it back after sufficient ticks have elapsed.

**Any implementation using busy-waiting will not get full points!**

## Task 3: Implement Priority Scheduler

- When threads are blocked and waiting for a lock, semaphore or a condition variable, the highest priority waiting thread should be woken up first.
- Implementation details:
  - Compare priorities of waiting threads when releasing locks, semaphores, condition variables.



# Priority Inversion

- Strict priority scheduling can lead to a phenomenon called “priority inversion”
- Supplemental reading:
  - What really happened to the pathfinder on Mars?
- Consider the following example where  
 $\text{prio}(H) > \text{prio}(M) > \text{prio}(L)$   
 $H$  needs a lock currently held by  $L$ , so  $H$  blocks  
 $M$  that was already on the ready list gets the processor before  $L$   
 $H$  indirectly waits for  $M$
- On Pathfinder, a watchdog timer noticed that  $H$  failed to run for some time, and would reset the system.

## Task 4: Implement Priority Donation

- When a high priority thread  $H$  waits on a lock held by a lower priority thread  $L$ , donate  $H$ 's priority to  $L$
- Recall the donation once  $L$  releases the lock.

Important:

- Remember to return  $L$  to previous priority once it releases the lock.
- Be sure to handle multiple donations (Multiple threads donating to a single thread)
- Be sure to handle nested donations ( $H$  waits on  $M$ , which waits on  $L$ , ...) up to 8 levels

# Synchronization

- Any synchronization problem can be easily solved by turning interrupts off: while interrupts are off, there is no concurrency, so there's no possibility for race conditions.
- **You should NOT do this, unless it is necessary!**
- Instead, use **semaphores**, **locks**, and **condition variables** to solve the bulk of your synchronization problems.
- **Exception:**  
**The only place you are allowed to turn interrupts off is**, when coordinating data shared between a kernel thread and interrupt handler. Because interrupt handlers can't sleep, they can't acquire locks.
- **Turning the interrupts off for synchronization between kernel threads (where it's not necessary) will lose points.**



## Task 5: Implement Advanced Scheduler

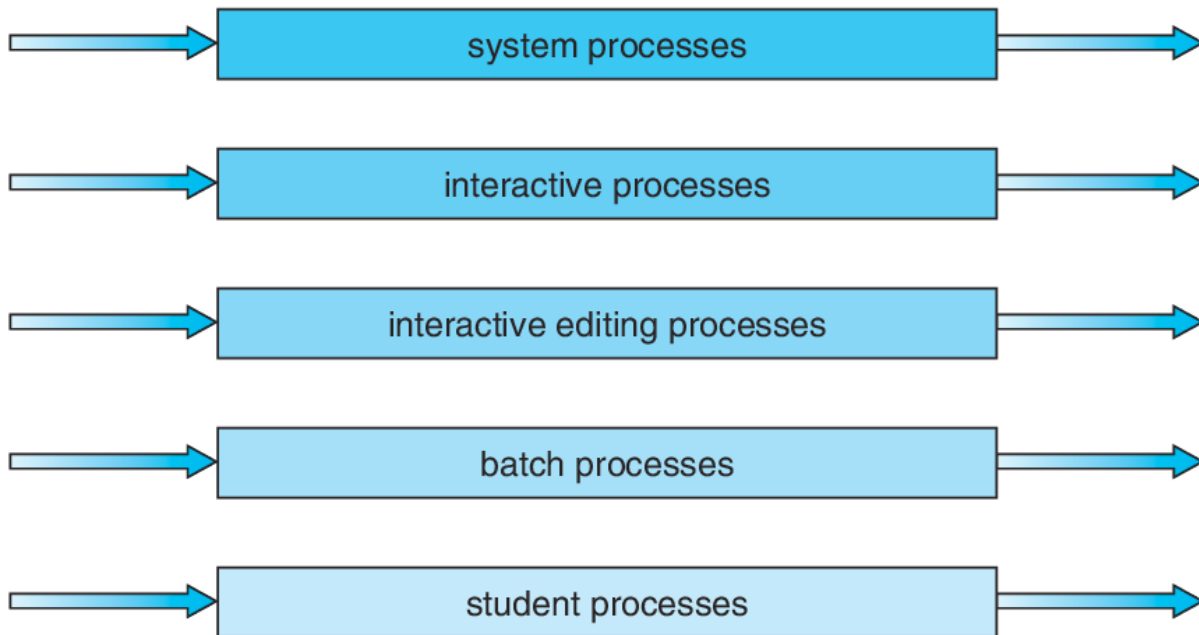
- Implement Multilevel Feedback Queue Scheduler
- Priority donation not needed in the advanced scheduler
  - Only one is active at a time
- Advanced scheduler must be chosen only if “-mlfqs” kernel option is specified

```
/* If false (default), use round-robin scheduler.  
   If true, use multi-level feedback queue scheduler.  
   Controlled by kernel command-line option "-o mlfqs". */  
bool thread_mlfqs;
```

- Read Appendix B - 4.4BSD Scheduler in Pintos manual for detailed information.
- Some of the parameters are real numbers and calculation involving them have to be simulated using integers.

# Multilevel Queue Scheduler

highest priority



lowest priority


# Multilevel Feedback Queue Scheduler

- A process can move between the various queues; aging can be implemented this way
- Multilevel Feedback Queue scheduler defined by the following parameters:
  - Number of queues
  - Scheduling algorithms for each queue
  - Method used to determine which queue a process will enter when that process needs service
  - Method used to determine when to upgrade a process
  - Method used to determine when to degrade a process

# Example of Multilevel Feedback Queue Scheduler

- Three queues:

- $Q_0$  - RR with  $q = 8$  ms
- $Q_1$  - RR with  $q = 16$  ms
- $Q_2$  - FCFS

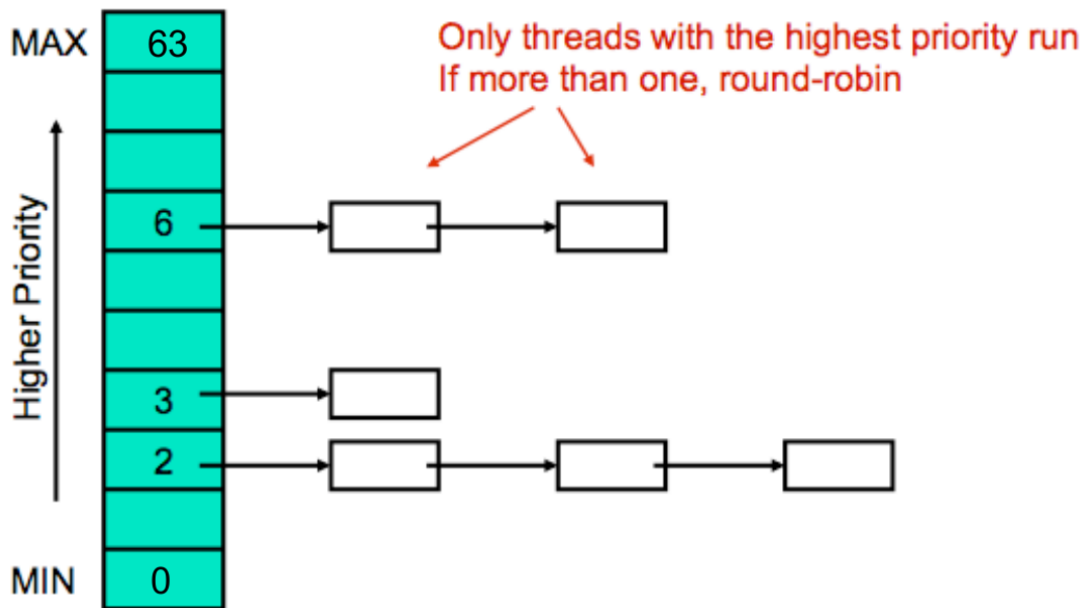


- Scheduling

- A new job enters queue  $Q_0$  which is served FCFS. When it gains CPU, job receives 8 milliseconds. If it does not finish in 8 milliseconds, job is moved to queue  $Q_1$
- At  $Q_1$  job is again served FCFS and receives 16 additional milliseconds. If it still does not complete, it is preempted and moved to queue  $Q_2$

# MLFQS: 4.4BSD Priority Based Scheduler

4.4BSD scheduler has 64 priorities and thus 64 ready queues, numbered 0 (PRI\_MIN) through 63 (PRI\_MAX).



# MLFQS: Calculating Priority

- NOTE: Lower numbers correspond to lower priorities in 4.4BSD, so that priority 0 is the lowest priority and priority 63 is the highest.
- Every 4 clock ticks, calculate:

$$\text{priority} = \text{PRI\_MAX} - (\text{recent\_cpu} / 4) - (\text{nice} * 2)$$

(rounded down to the nearest integer)

- It gives a thread that has received CPU time recently lower priority for being reassigned the CPU the next time the scheduler runs. (Aging)

# MLFQS: “Nice” Value

How “*nice*” the thread should be to other threads.

- A nice of zero does not affect thread priority.
- A positive nice, to the **maximum of +20**, decreases the priority of a thread and causes it to give up some CPU time it would otherwise receive.
- A negative nice, to the **minimum of -20**, tends to take away CPU time from other threads.

[Read Pintos manual Section “B. 4.4BSD Scheduler” for details](#)

# MLFQS: Calculating “*recent\_cpu*”

- An array of  $n$  elements to track the CPU time received in each of the last  $n$  seconds requires  $O(n)$  space per thread and  $O(n)$  time per calculation of a new weighted average.
- Instead, we use an exponentially weighted moving average:
  - **`recent_cpu(0) = 0`** for initial thread, and parent thread's value for other threads.
  - At each timer interrupt, **`recent_cpu`** incremented by 1 for the running thread.
  - And once per second, for each thread:

$$a = (2 * \text{load\_avg}) / (2 * \text{load\_avg} + 1)$$
$$\text{recent\_cpu}(t) = a * \text{recent\_cpu}(t - 1) + \text{nice}$$



# MLFQS: Calculating “*load\_avg*”

- Estimates the average number of threads ready to run over the past minute.
- Like `recent_cpu`, it is an exponentially weighted moving average.
- Unlike `priority` and `recent_cpu`, `load_avg` is system-wide, not thread-specific.
- At system boot, it is **initialized to 0**. Once per second thereafter, it is updated according to the following formula:

$$\text{load\_avg}(t) = (59 / 60) * \text{load\_avg}(t - 1) + (1 / 60) * \text{ready\_threads}$$

- *ready\_threads*: number of threads that are either running or ready to run at the time of update

# Functions to implement

Skeletons of these functions are provided in `src/threads/threads.c`

- `int thread_get_nice (void)`
- `void thread_set_nice (int new_nice)`
- `void thread_set_priority (int new_priority)`
- `int thread_get_priority (void)`
- `int thread_get_recent_cpu (void)`
- `int thread_get_load-avg (void)`



# Suggested Order of Implementation

## Phase-1

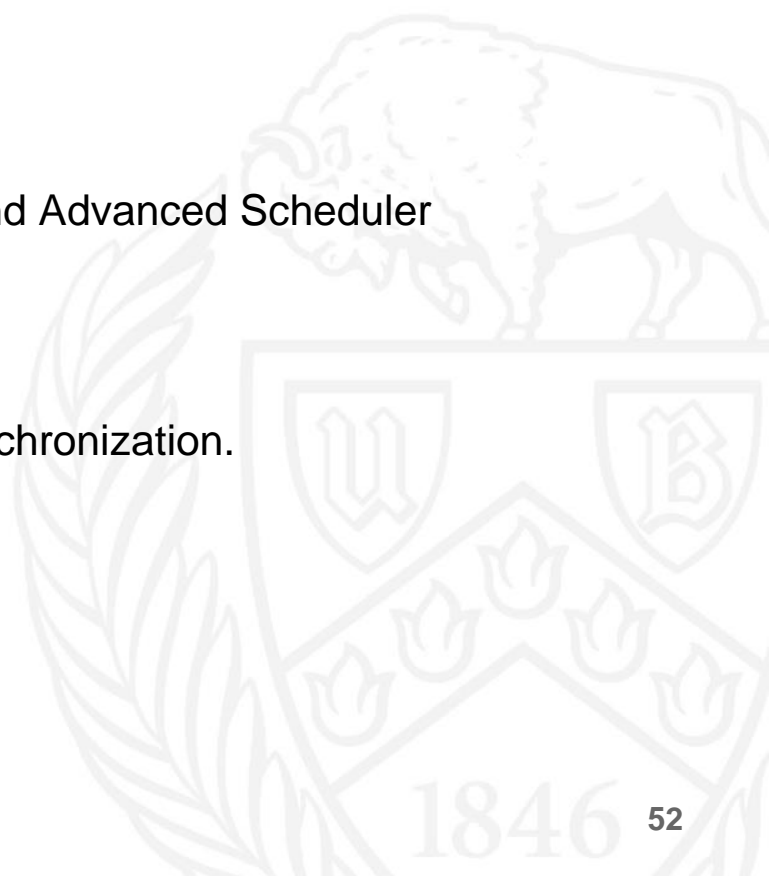
- Priority Scheduler
  - Needed for implementing Priority Donation and Advanced Scheduler

## Phase-2

- Alarm Clock
  - Needs more understanding of Pintos and synchronization.
  - Other parts do not depend on this

## Phase-3

- Priority Scheduler (cont.)
- Priority Donation
- Advanced Scheduler



# Debugging your Code

- printf, ASSERT, backtraces, gdb
- Running pintos under gdb
  - Invoke pintos with the gdb option (Note the spaces and hyphens).
    - `pintos --gdb -- run testname` Do not copy command from PDF
  - On another terminal from build/ directory, invoke gdb
    - `pintos-gdb kernel.o`
  - Issue the command
    - `debugpintos`
  - All the usual gdb commands can be used: step, next, print, continue, break, clear, etc.
  - Psst... Use the pintos debugging macros described in manual (e.g. `dumplist`) <sup>53</sup>

# How Much Code?

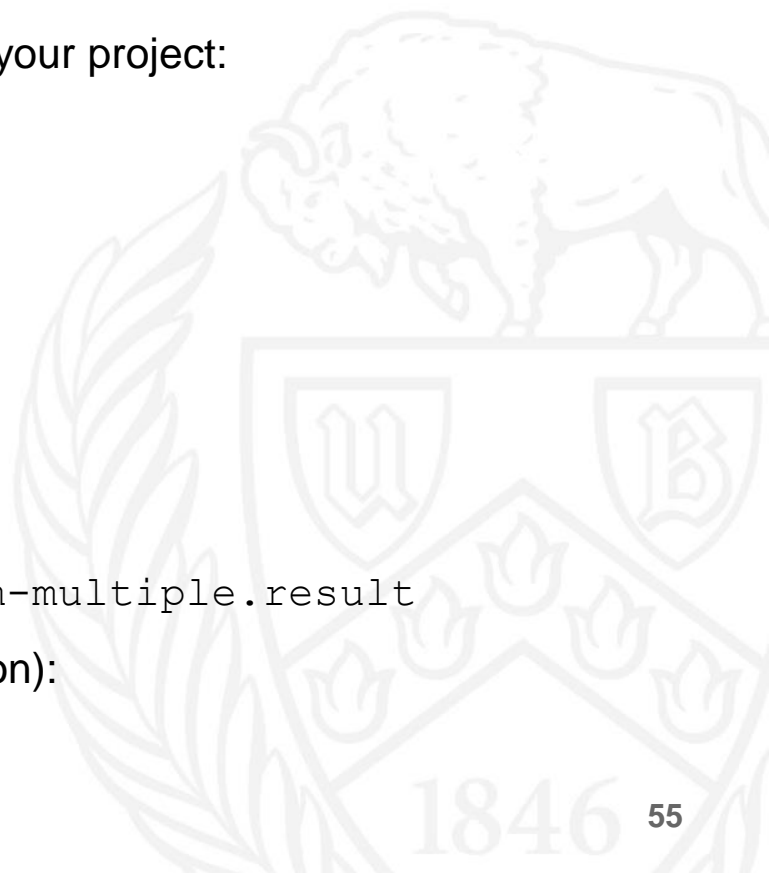
devices/timer.c		42	+++++-
threads/fixed-point.h		120	+++++
threads/synch.c		88	+++++
threads/thread.c		196	+++++-----
threads/thread.h		23	+++

5 files changed, 440 insertions(+), 29 deletions(-)

- This reference solution represents just one possible solution.

## Step 5: Testing

- Pintos provides a very systematic testing suite for your project:
  - Compile:
    - `make clean`
    - `make`
  - Run all tests (pass/fail):
    - `make check`
  - Run individual tests:
    - `make build/tests/threads/alarm-multiple.result`
  - Run the grading script (gives useful information):
    - `make grade`



## Step 5: Testing

- **make check**

```
pass tests/threads/alarm-single
pass tests/threads/alarm-multiple
pass tests/threads/alarm-simultaneous
FAIL tests/threads/alarm-priority
pass tests/threads/alarm-zero
pass tests/threads/alarm-negative
FAIL tests/threads/priority-change
FAIL tests/threads/priority-donate-one
FAIL tests/threads/priority-donate-multiple
FAIL tests/threads/priority-donate-multiple2
FAIL tests/threads/priority-donate-nest
FAIL tests/threads/priority-donate-sema
FAIL tests/threads/priority-donate-lower
FAIL tests/threads/priority-fifo
FAIL tests/threads/priority-preempt
FAIL tests/threads/priority-sema
FAIL tests/threads/priority-condvar
FAIL tests/threads/priority-donate-chain
```

```
FAIL tests/threads/mlfqs-load-1
FAIL tests/threads/mlfqs-load-60
FAIL tests/threads/mlfqs-load-avg
FAIL tests/threads/mlfqs-recent-1
pass tests/threads/mlfqs-fair-2
pass tests/threads/mlfqs-fair-20
FAIL tests/threads/mlfqs-nice-2
FAIL tests/threads/mlfqs-nice-10
FAIL tests/threads/mlfqs-block
20 of 27 tests failed.
```

# Grading

- **make grade**
- `src/threads/build/grade:`

**TOTAL TESTING SCORE: 24.2%**

-----

## SUMMARY BY TEST SET

Test Set	Pts	Max	%	Ttl	%	Max
tests/threads/Rubric.alarm	14/	18	15.6%/	20.0%		
tests/threads/Rubric.priority	0/	38	0.0%/	40.0%		
tests/threads/Rubric.mlfqs	8/	37	8.6%/	40.0%		
Total			24.2%/	100.0%		



# Grading - Alarm Clock: 14/18 pts

```
4/ 4 tests/threads/alarm-single
4/ 4 tests/threads/alarm-multiple
4/ 4 tests/threads/alarm-simultaneous
** 0/ 4 tests/threads/alarm-priority

1/ 1 tests/threads/alarm-zero
1/ 1 tests/threads/alarm-negative
```

Phase-2

- Section summary.

5/ 6 tests passed

14/ 18 points subtotal

All tests in Phase-1 are included in Phase-2.

- If alarm clock implementation is based on “busy-waiting”, or
- If interrupts are turned off excessively for synchronization (between kernel threads)

➤ **You will lose points**

# Grading - Priority Scheduler: 0/38 pts

```
** 0/ 3 tests/threads/priority-change  
** 0/ 3 tests/threads/priority-preempt
```

Phase-1

```
** 0/ 3 tests/threads/priority-fifo  
** 0/ 3 tests/threads/priority-sema  
** 0/ 3 tests/threads/priority-condvar
```

```
** 0/ 3 tests/threads/priority-donate-one  
** 0/ 3 tests/threads/priority-donate-multiple  
** 0/ 3 tests/threads/priority-donate-multiple2  
** 0/ 3 tests/threads/priority-donate-nest  
** 0/ 5 tests/threads/priority-donate-chain  
** 0/ 3 tests/threads/priority-donate-sema  
** 0/ 3 tests/threads/priority-donate-lower
```

Phase-3

- Section summary.
  - 0/ 12 tests passed
  - 0/ 38 points subtotal

# Grading - MLFQ Scheduler: 8/37 pts

```
** 0/ 5 tests/threads/mlfqs-load-1
** 0/ 5 tests/threads/mlfqs-load-60
** 0/ 3 tests/threads/mlfqs-load-avg

** 0/ 5 tests/threads/mlfqs-recent-1

    5/ 5 tests/threads/mlfqs-fair-2
    3/ 3 tests/threads/mlfqs-fair-20

** 0/ 4 tests/threads/mlfqs-nice-2
** 0/ 2 tests/threads/mlfqs-nice-10

** 0/ 5 tests/threads/mlfqs-block

- Section summary.
    2/ 9 tests passed
    8/ 37 points subtotal
```

Phase-3



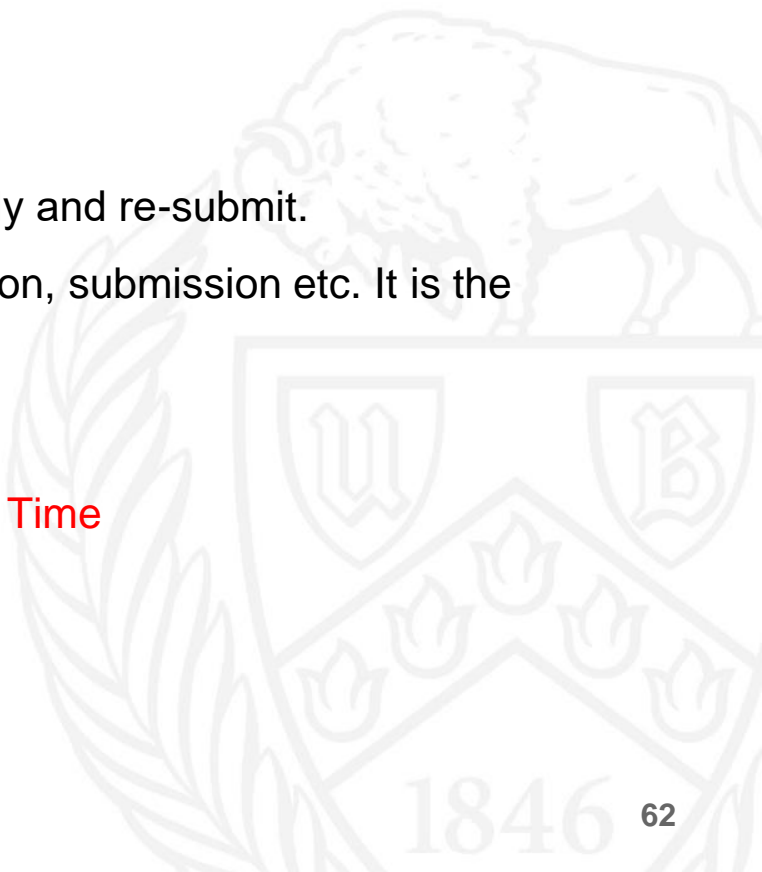
All tests in “make check” are included in Phase-3.

# Grading

- Your “source code” consists of weighted average of the three phases.  
(10%, 25%, 65% and Due dates: Sep 17<sup>th</sup>, Oct 1<sup>st</sup>, Oct 15<sup>th</sup>)
- Your “design document” is 10% of your Project-1 grade, and “source code” is 90%
- The points are weighted. `make grade` will give you a score out of 100% for your code, based on the passed tests and their weight.
- You can consider that for summary of your tests. Our grading scheme is different, as our requirements for phase-1, phase-2, phase-3 are different.
- Check autograder submission score to be consistent with what you get on your VM.
  - Reminder: points will be taken off for busy-waiting and/or turning interrupts off unnecessarily.


# Submission

- Submission will be via AutoLab autograder.
  - The instructions will be posted on Piazza.
  - You'll have unlimited submissions, submit early and re-submit.
  - Every phase needs registration, group formation, submission etc. It is the responsibility of both members to ensure this.
- Due days and times are Fridays 11:59 PM Eastern Time
- Refer to **LATE SUBMISSION** policy.



# Assignments

- If not yet, submit/join your groups on github classroom.
- If you do not have a group, post a private note on Piazza.
- Understand the initial pintos code in the files related to Project-1.
  - thread.c , thread.h
  - timer.c , timer.h
  - init.c , init.h
  - interrupt.c , interrupt.h
  - synch.c , synch.h
  - list.c , list.h
- Get started!!!
- Reading: Chapter 4 (Concurrency) from zyBooks



# Summary

- Pintos projects:
  - Project-1: Threads
    - Step 1: Preparation
    - Step 2: Setting Up Pintos
    - Step 3: Design Document
    - Step 4: Implementation
    - Step 5: Testing
  - Project-2: User Programs
  - Project-3: Virtual Memory
  - Project-4: File Systems



# Acknowledgements

- “Operating Systems Concepts” book and supplementary material by A. Silberschatz, P. Galvin and G. Gagne
- “Operating Systems: Internals and Design Principles” book and supplementary material by W. Stallings
- “Modern Operating Systems” book and supplementary material by A. Tanenbaum
- R. Doursat and M. Yuksel from University of Nevada, Reno
- T. Kosar and K. Dantu from University at Buffalo
- Pintos Manual
- Pintos Notes and Slides by A. He (Stanford), A. Romano (Stanford), J. Sundararaman & X. Liu(Virginia Tech), J. Kim (SKKU)